

What using Claude costs.

Use is the material the model processes — measured in tokens.

Use is the material the model processes

An LLM works through everything you give it and everything it generates back.



INPUTS

What you give it

Your prompt, any files, and the running conversation.



OUTPUTS

What it generates

Everything Claude writes back to you.


■ **Both sides count.** The more material in and out, the more use.

A token is a small chunk of text.

LLMs break all that material — inputs and outputs — into tokens, then process them one at a time.

Cost = tokens × the model you pick

The same task costs more on a more capable model. Each tier up roughly doubles or triples your per-task spend.



Rule of thumb: start with Sonnet; reach for Opus only when it falls short.



PER-MODEL RATES

Token costs by model: the Cornell AI API Gateway page (Confluence). You can also check external sites like pricepertoken.com.

Mostly tokens, plus a little platform overhead

Usage cost is primarily tokens, with small charges for running Cornell's AI Platform.



TOKENS

The primary cost

Priced per model — see the gateway page for rates.



GATEWAY

\$0.002 / request

\$2 per 1,000 requests, per LiteLLM licensing.



WEB SEARCH

~\$0.05 / search

Per Gemini search, n8n-driven (being finalized).

■ **All included in your cap.** Tokens plus this nominal overhead count toward your monthly limit.

Your administrator sets a monthly cap

Usage is capped per person, so spend stays predictable.



WITHIN YOUR CAP

- Work normally
- Nothing to manage
- e.g. \$20 / month per person



IF YOU EXCEED IT

- You'll get an error message
- Contact your administrator
- They can raise your cap

Caps keep spend predictable — **no surprise bills.**