


COST & TOKENS

Use is billed by tokens.

Claude bills by the **token** — the small chunks of text the model reads and writes as it builds an answer one token at a time. You pay for tokens *in* (your prompt, any files, and the running conversation) plus tokens *out* (Claude's reply). Cost flexes with the work: a quick question is a few tokens; a long thread with big attachments is many.



There's no flat per-prompt or per-session price. Two things set the bill: **how many tokens** move through, and **which model** you pick. Usage is billed to your unit through the Cornell AI Platform, and most subscription features are included.

Pick the model to fit the job

Haiku · \$

Quick lookups, classification, short rewrites. Fastest and lowest cost.

★ Sonnet · \$\$

The workhorse — drafting, summarizing, analysis. Start here for almost any task.

Opus · \$\$\$

Long documents and hard reasoning. Reach for it when Sonnet falls short.

\$ / \$\$ / \$\$\$ is relative cost — each tier up roughly doubles or triples your per-task spend.

Keep costs in hand

- **Default to Sonnet;** use Haiku for small tasks and move to Opus only when needed.
- **Keep context lean** — start a fresh chat for a new task and attach only the files Claude needs.
- **Match the model to the job** rather than running everything on the top tier.

For current per-token rates by model, the most popular tracker is **Price Per Token**.